

This study note is heavily based on [this EMNLP 2018 tutorial](#). It covers the math behind VAEs.

Bridging Latent Variable Models and Deep Learning

Deep Learning

- Broadly construed, deep learning is a toolbox for learning rich representations of data through numerical optimization.
- Deep learning makes it possible to parameterize conditional likelihoods with powerful function approximations.

Latent Variable Models

- LVMs make it easy to explicitly specify model constraints through conditional independence properties.
- LVMs objectives often complicate backpropagation by introducing points of non-differentiability.

Targeted Issue and Main Focuses

How to combine the complementary strengths of the both worlds, and address the issues?

Variational inference a key technique for performing approximate posterior inference

The main focus of the [tutorial](#) is on training an inference network (deep inference) to output (*latent variable inference*) the parameters of an approximate posterior distribution given the set of variables to be conditioned upon. Also, it focuses on **learning LVMs whose joint distribution can be expressed as a directed graphical model** (DGM), which is done through variational inference.

Learning and Inference

We are interested in two things after the model is defined:

- Learning model parameter θ
- Performing *inference* after learning θ (computing the posterior distribution $p(z|x; \theta)$, or approximated, over the latent variables, given some data x)

These two tasks are intimately connected because learning often uses inference as a subroutine.

Learning often involves alternatively inferring likely z values, and optimizing the model assuming these inferred z 's. The dominating approach to train a latent variable model is through maximizing likelihood.

Tractable case

Assuming tractable log marginal likelihood, i.e.

$$\log p(x; \theta) = \log \sum_z p(x, z; \theta) = \log \sum_z p(z|x; \theta) p(x; \theta)$$

is tractable to evaluate (which is equivalent to assuming posterior inference to be tractable).

- Categorical LVMs where K is not too big
- HMMs where dynamic programs allow us to efficiently sum over all the z assignments

Using maximum likelihood training, learning θ then corresponds to solving the following problem:

$$\operatorname{argmax}_{\theta} \sum_{n=1}^N \log p(x^{(n)}; \theta)$$

where we have assumed N examples in our training set.

Directly Optimizing with Gradient Descent

$$L(\theta) = \log p(x^{(1:N)}; \theta) = \sum_{n=1}^N \log p(x^{(n)}; \theta) = \sum_{n=1}^N \log \sum_z p(x^{(n)}, z; \theta)$$

where the gradient is

$$\nabla_{\theta} L(\theta) = \sum_{n=1}^N \mathbb{E}_{p(z|x^{(n)}; \theta)} [\nabla_{\theta} p(x^{(n)}, z; \theta)]$$

which is the same form of the M-step in EM algorithm (to be confirmed). Note that the gradient expression involves an expectation over the posterior $p(z|x^{(n)}; \theta)$, and *is an example of how inference is used as a subroutine in learning*. By gradient descent

$$\theta^{(i+1)} = \theta^{(i)} + \eta \nabla_{\theta} L(\theta^{(i)}).$$

Expectation Maximization (EM) Algorithm

EM is an iterative method for learning LVMs with tractable posterior inference. It maximizes a lower bound on the log marginal likelihood at each iteration. Given randomly-initialized starting parameters $\theta^{(0)}$, the algorithm updates the parameters via the following alternating procedure:

1. E-step: Derive the posterior under current parameters $\theta^{(i)}$, i.e.,

$$p(z|x^{(n)}; \theta^{(i)}) = \frac{p(x^{(n)}, z; \theta^{(i)})}{p(x^{(n)})}$$

for all $n = 1, \dots, N$.

2. M-step: Define the *expected complete data likelihood* as

$$Q(\theta, \theta^{(i)}) = \sum_{n=1}^N \mathbb{E}_{p(z|x^{(n)}; \theta^{(i)})} [\log p(x^{(n)}, z; \theta)]$$

Maximize this w.r.t. θ while holding $\theta^{(i)}$ fixed

$$\theta^{(i+1)} = \operatorname{argmax}_{\theta} Q(\theta, \theta^{(i)})$$

It can be shown that EM improves the log marginal likelihood at each iteration, i.e.

$$\sum_{n=1}^N \log p(x^{(n)}; \theta^{(i+1)}) \geq \sum_{n=1}^N \log p(x^{(n)}; \theta^{(i)})$$

In some cases, there is an exact solution to M-step (e.g., GMMs); otherwise, one can use gradient descent and the expression is the same as directly optimizing the log likelihood. This variant of EM (no exact M-step solution) is referred to as *generalized EM*.

Intractable case

What if calculation of posterior inference or log marginal likelihood is intractable? Variational inference is a technique for approximating an intractable posterior distribution $p(z|x; \theta)$ with a tractable surrogate. In the context of learning the parameters of LVMS, VI can be used in optimizing a lower bound on the log marginal likelihood that involves only an approximate posterior over latent variables, rather than the exact posteriors.

$$\begin{aligned} \log p(x; \theta) &= \int q(z; \lambda) \log p(x; \theta) dz \\ &= \int q(z; \lambda) \log \frac{p(x, z; \theta)}{p(z|x; \theta)} dz \\ &= \int q(z; \lambda) \log \frac{p(x, z; \theta) q(z; \lambda)}{q(z; \lambda) p(z|x; \lambda)} dz \\ &= \int q(z; \lambda) \log \frac{p(x, z; \theta)}{q(z; \lambda)} dz + \int q(z; \lambda) \log \frac{q(z; \lambda)}{p(z|x; \lambda)} dz \\ &= \mathbb{E}_{q(z; \lambda)} \log \frac{p(x, z; \theta)}{q(z; \lambda)} + \text{KL}[q(z; \lambda) || p(z|x; \lambda)] \\ &= \text{ELBO}(\theta, \lambda; x) + \text{KL}[q(z; \lambda) || p(z|x; \lambda)] \\ &\geq \text{ELBO}(\theta, \lambda; x) \end{aligned}$$

Given N data, the ELBO over the entire dataset is given by the sum of individual ELBOs ($x^{(n)}$ are assumed to be drawn i.i.d),

$$\text{ELBO}(\theta, \lambda; x^{(1:N)}) = \sum_{n=1}^N \mathbb{E}_{q(z; \lambda^{(n)})} \left[\log \frac{p(x^{(n)}, z; \theta)}{q(z; \lambda^{(n)})} \right] \leq \log p(x^{(1:N)}; \theta)$$

Note that $\lambda = [\lambda^{(1)}, \dots, \lambda^{(n)}]$ (i.e., we have $\lambda^{(n)}$ for each data point $x^{(n)}$), which will be further approximated in the context of variational autoencoders (VAEs) with amortized variational inference. It is this aggregate ELBO that we wish to maximize w.r.t. θ and λ to train our model.

Maximizing the aggregate ELBO

coordinate ascent (variational EM)

1. Variational E-step: Maximize the ELBO for each $x^{(n)}$ holding $\theta^{(i)}$ fixed

$$\begin{aligned}
\lambda^{(n)} &= \operatorname{argmax}_{\lambda} \operatorname{ELBO}(\theta^{(i)}, \lambda; x^{(n)}) \\
&= \operatorname{argmax}_{\lambda} \mathbb{E}_{q(z; \lambda)} \left[\log \frac{p(z|x^{(n)}; \theta) p(x^{(n)}; \theta)}{q(z; \lambda)} \right] \\
&= \operatorname{argmin}_{\lambda} \mathbb{E}_{q(z; \lambda)} \left[\log \frac{q(z; \lambda)}{p(z|x^{(n)}; \theta)} \right] \\
&= \operatorname{argmin}_{\lambda} \operatorname{KL}[q(z; \lambda) || p(z|x^{(n)}; \theta)],
\end{aligned}$$

where the third equality holds since $\log p(x; \theta^{(i)})$ is a constant w.r.t. $\lambda^{(n)}$'s.

2. Variational M-step: Maximize the aggregated ELBO w.r.t. θ holding the $\lambda^{(n)}$'s fixed

$$\theta^{(i+1)} = \operatorname{argmax}_{\theta} \sum_{n=1}^N \operatorname{ELBO}(\theta, \lambda^{(n)}; x^{(n)}) = \operatorname{argmax}_{\theta} \sum_{q(z; \lambda^{(n)})} \mathbb{E}_{q(z; \lambda)} [\log p(x^{(n)}, z; \theta)],$$

where the second equality holds since the term $\mathbb{E}_{q(z; \lambda)} [-\log q(z; \lambda^{(n)})]$ is constant w.r.t. θ .

This is known as **variational expectation maximization**, because the E-step is replaced with variational inference which finds the best variational approximation to the true posterior. The M-step maximizes the expected complete data likelihood where the expectation is taken w.r.t. the variational posterior.

If for each data $x^{(n)}$ there exists $\lambda^{(n)}$ such that $q(z; \lambda^{(n)}) = p(z|x^{(n)}; \theta)$, we say that *the variational family is flexible enough to include the true posterior*, and **it reduces to the classic EM algorithm**. However, we are interested in finding a flexible variational family that allows for tractable optimization since we have assumed the exact posterior inference is intractable.